

Perfilação

Computação de Alto Desempenho

Lucas Ribeiro Ikuhara - DRE 119019172

Considerações iniciais

Ao analisarmos o código, pode-se perceber que se trata de uma aproximação de uma equação de laplace, para que possa ser resolvida usando o método de relaxação. Essa técnica é comumente usada para resolver numericamente equações elípticas. A aplicação analisada é de caráter científico, e pode-se imaginar que pode demorar muito dependendo das entradas. Para melhorar a performance da implementação, pode-se usar um perfilador, para descobrirmos quais partes do código gastam mais tempo para serem executadas no total (hotspots), para dirigirmos nossos esforços de otimização.

Passos para perfilação

Para conseguirmos usar o gprof, é necessário que o código seja compilado com os flags `-pg`, e que o executável gerado seja executado pelo menos uma vez. Quando a execução é finalizada, é gerado um arquivo `gmon.out`. Nesse ponto, o perfilador pode ser chamado.

Utilizando o g++ e gprof, o processo de compilação e perfilação pode ser exemplificado assim:

```
$ g++ laplace.cxx -pg -o saida.out
$ ./saida.out
Enter nx n_iter eps --> 500 100 1e-6
$ gprof saida.out
```

Resultados da perfilação

O perfilador gera um relatório detalhado com explicações adicionais, que serão coladas integralmente no final do relatório. Os resultados mais relevantes, no entanto, foram compilados em uma tabela, com o nome das funções, quantidade de chamadas e tempo gasto em cada uma.

Flat profile:

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	ms/call	ms/call	name
68.44	0.15	0.15	100	1.51	2.16	
LaplaceSolver::timeStep(double)						
29.66	0.22	0.07	24800400	0.00	0.00	SQR(double const&)
2.28	0.22	0.01	2	2.51	2.51	seconds()
0.00	0.22	0.00	2000	0.00	0.00	BC(double, double)

0.00	0.22	0.00	1	0.00	0.00
_GLOBAL__sub_I_ZN4GridC2Eii					
0.00	0.22	0.00	1	0.00	0.00
__static_initialization_and_destruction_0(int, int)					
0.00	0.22	0.00	1	0.00	0.00
LaplaceSolver::initialize()					
0.00	0.22	0.00	1	0.00	215.81
LaplaceSolver::solve(int, double)					
0.00	0.22	0.00	1	0.00	0.00
LaplaceSolver::LaplaceSolver(Grid*)					
0.00	0.22	0.00	1	0.00	0.00
LaplaceSolver::~~LaplaceSolver()					
0.00	0.22	0.00	1	0.00	0.00
Grid::setBCFunc(double (*) (double, double))					
0.00	0.22	0.00	1	0.00	0.00
Grid::Grid(int, int)					

Ao analisar os resultados explicitados, podemos ver que os dois hotspots principais são as funções `LaplaceSolver::timeStep` e `SQR`, que deverão ser nosso foco ao tentar otimizar o código.

Ao analisarmos o código fonte dessas duas funções, é fácil perceber que `Real LaplaceSolver :: timeStep` não segue boas práticas de CAD. Entretanto, a função `SQR` já se encontra relativamente otimizada e seria difícil conseguir ganhos de performance alterando ela.

- Real LaplaceSolver :: timeStep

```
Real LaplaceSolver :: timeStep(const Real dt)
{
    Real dx2 = g->dx*g->dx;
    Real dy2 = g->dy*g->dy;
    Real tmp;
    Real err = 0.0;
    int nx = g->nx;
    int ny = g->ny;
    Real **u = g->u;

    for (int i=1; i<nx-1; ++i) {
        for (int j=1; j<ny-1; ++j) {
            tmp = u[i][j];
            u[i][j] = ((u[i-1][j] + u[i+1][j])*dy2 +
                       (u[i][j-1] + u[i][j+1])*dx2)*0.5/(dx2 + dy2);
            err += SQR(u[i][j] - tmp);
        }
    }
}
```

```

    return sqrt(err);
}

```

Nessa função, podemos imediatamente observar que os loops realizam operações para computar sua condição de parada, que poderia ter sido feitas fora de sua declaração. Além disso, são realizadas divisões custosas no loop que poderiam ter sido evitadas.

Otimizando o código

Resolvendo os problemas analisados acima, tentaremos melhorar a performance do programa. O tempo de execução inicial reportado pelo programa é 0.446201 segundos.

Como anteriormente citado, deve-se evitar operações de divisão custosas dentro de loops; Dessa forma, em `LaplaceSolver :: timeStep`, o trecho `u[i][j] = ((u[i-1][j] + u[i+1][j])*dy2 + (u[i][j-1] + u[i][j+1])*dx2);` foi alterado, e partes dele foram calculadas antes de começar a iteração. Visando reduzindo a quantidade de multiplicações e divisões por iteração, temos:

```

Real k = 0.5/(dx2 + dy2);
dx2 = dx2 * k;
dy2 = dy2 * k;

for (int i=1; i<nx-1; ++i) {
    for (int j=1; j<ny-1; ++j) {
        tmp = u[i][j];
        u[i][j] = (u[i-1][j] + u[i+1][j])*dy2 +
                  (u[i][j-1] + u[i][j+1])*dx2;
        err += SQR(u[i][j] - tmp);
    }
}

```

Após essa alteração, o tempo de execução foi reduzido de 0.446201 segundos para 0.417577 segundos.

Foram testadas mais duas alterações, mas nenhuma delas melhorou a performance. A primeira foi tentar alterar a função `SQR`, alterando o trecho `(x*x)` por `pow(x, 2)`. A ideia por trás dessa mudança é que possivelmente a implementação de exponenciação da biblioteca `cmath` seria mais rápida do que a operação inicial. essa hipótese foi rapidamente refutada, uma vez que a mudança dobrou o tempo de execução do programa.

A outra alteração testada foi substituir a condição de parada do loop `for (int i=1; i<nx-1; ++i)` para calcular `nx-1` antes de sua declaração. Esa alteração não aumentou nem diminuiu apreciavelmene o tempo de execução. Como em média testando algumas vezes seguidas o tempo aumentou, mesmo que por muito pouco, a alteração foi descartada.

Além disso, cosiderou-se a possibilidade de usar as diretivas `#pragma` para tentar forçar a vetorização do loop. Entretanto, olhando mais atenciosamente o código, pode-se perceber que as iterações são dependentes uma da outra.

Conclusões finais

É evidente que perfiladores são ferramentas muito poderosas quando se busca extrair mais performance. Não só porquê nos dão ferramentas robustas para conduzir testes relacionados a performance, mas porquê tornam o processo como um todo mais objetivo e científico, possibilitando que esforços sejam concentrados nas partes mais importantes do código: os hotspots. O ganho total de performance foi de aproximadamente 6.41%, que adimitidamente modesto, não é desprezível, e

Saída do perfilador completa

Flat profile:

Each sample counts as 0.01 seconds.

%	cumulative	self	self	total	
time	seconds	seconds	calls	ms/call	ms/call name
68.44	0.15	0.15	100	1.51	2.16
LaplaceSolver::timeStep(double)					
29.66	0.22	0.07	24800400	0.00	0.00 SQR(double const&)
2.28	0.22	0.01	2	2.51	2.51 seconds()
0.00	0.22	0.00	2000	0.00	0.00 BC(double, double)
0.00	0.22	0.00	1	0.00	0.00
_GLOBAL__sub_I_ZN4GridC2Eii					
0.00	0.22	0.00	1	0.00	0.00
__static_initialization_and_destruction_0(int, int)					
0.00	0.22	0.00	1	0.00	0.00
LaplaceSolver::initialize()					
0.00	0.22	0.00	1	0.00	215.81
LaplaceSolver::solve(int, double)					
0.00	0.22	0.00	1	0.00	0.00
LaplaceSolver::LaplaceSolver(Grid*)					
0.00	0.22	0.00	1	0.00	0.00
LaplaceSolver::~~LaplaceSolver()					
0.00	0.22	0.00	1	0.00	0.00 Grid::setBCFunc(double
(*) (double, double))					
0.00	0.22	0.00	1	0.00	0.00 Grid::Grid(int, int)
%					
time	the percentage of the total running time of the program used by this function.				

cumulative a running sum of the number of seconds accounted
seconds for by this function and those listed above it.

self the number of seconds accounted for by this
seconds function alone. This is the major sort for this
listing.

calls the number of times this function was invoked, if
this function is profiled, else blank.

self the average number of milliseconds spent in this
ms/call function per call, if this function is profiled,
else blank.

total the average number of milliseconds spent in this
ms/call function and its descendents per call, if this
function is profiled, else blank.

name the name of the function. This is the minor sort
for this listing. The index shows the location of
the function in the gprof listing. If the index is
in parenthesis it shows where it would appear in
the gprof listing if it were to be printed.

Copyright (C) 2012-2020 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification,
are permitted in any medium without royalty provided the copyright
notice and this notice are preserved.

Call graph (explanation follows)

granularity: each sample hit covers 2 byte(s) for 4.53% of 0.22 seconds

index	% time	self	children	called	name
					<spontaneous>
[1]	100.0	0.00	0.22		main [1]
		0.00	0.22	1/1	LaplaceSolver::solve(int, double) [3]
		0.01	0.00	2/2	seconds() [5]

```

0.00 0.00 1/1 Grid::Grid(int, int) [19]
0.00 0.00 1/1 Grid::setBCFunc(double (*)
(double, double)) [18]
0.00 0.00 1/1
LaplaceSolver::LaplaceSolver(Grid*) [16]
0.00 0.00 1/1
LaplaceSolver::~~LaplaceSolver() [17]
-----
0.15 0.07 100/100 LaplaceSolver::solve(int,
double) [3]
[2] 97.7 0.15 0.07 100 LaplaceSolver::timeStep(double)
[2]
0.07 0.00 24800400/24800400 SQR(double const&) [4]
-----
0.00 0.22 1/1 main [1]
[3] 97.7 0.00 0.22 1 LaplaceSolver::solve(int,
double) [3]
0.15 0.07 100/100
LaplaceSolver::timeStep(double) [2]
-----
0.07 0.00 24800400/24800400
LaplaceSolver::timeStep(double) [2]
[4] 29.5 0.07 0.00 24800400 SQR(double const&) [4]
-----
0.01 0.00 2/2 main [1]
[5] 2.3 0.01 0.00 2 seconds() [5]
-----
0.00 0.00 2000/2000 Grid::setBCFunc(double (*)
(double, double)) [18]
[12] 0.0 0.00 0.00 2000 BC(double, double) [12]
-----
0.00 0.00 1/1 __libc_csu_init [24]
[13] 0.0 0.00 0.00 1 _GLOBAL__sub_I__ZN4GridC2Eii
[13]
0.00 0.00 1/1
__static_initialization_and_destruction_0(int, int) [14]
-----
0.00 0.00 1/1
_GLOBAL__sub_I__ZN4GridC2Eii [13]
[14] 0.0 0.00 0.00 1
__static_initialization_and_destruction_0(int, int) [14]
-----
0.00 0.00 1/1

```

```

LaplaceSolver::LaplaceSolver(Grid*) [16]
[15]      0.0      0.00      0.00      1      LaplaceSolver::initialize()
[15]
-----
      0.00      0.00      1/1      main [1]
[16]      0.0      0.00      0.00      1
LaplaceSolver::LaplaceSolver(Grid*) [16]
      0.00      0.00      1/1      LaplaceSolver::initialize()
[15]
-----
      0.00      0.00      1/1      main [1]
[17]      0.0      0.00      0.00      1      LaplaceSolver::~~LaplaceSolver()
[17]
-----
      0.00      0.00      1/1      main [1]
[18]      0.0      0.00      0.00      1      Grid::setBCFunc(double (*
(double, double)) [18]
      0.00      0.00      2000/2000      BC(double, double) [12]
-----
      0.00      0.00      1/1      main [1]
[19]      0.0      0.00      0.00      1      Grid::Grid(int, int) [19]
-----

```

This table describes the call tree of the program, and was sorted by the total amount of time spent in each function and its children.

Each entry in this table consists of several lines. The line with the index number at the left hand margin lists the current function.

The lines above it list the functions that called this function, and the lines below it list the functions this one called.

This line lists:

index	A unique number given to each element of the table. Index numbers are sorted numerically. The index number is printed next to every function name so it is easier to look up where the function is in the table.
% time	This is the percentage of the 'total' time that was spent in this function and its children. Note that due to different viewpoints, functions excluded by options, etc, these numbers will NOT add up to 100%.
self	This is the total amount of time spent in this function.

children	This is the total amount of time propagated into this function by its children.
called	This is the number of times the function was called. If the function called itself recursively, the number only includes non-recursive calls, and is followed by a '+' and the number of recursive calls.
name	The name of the current function. The index number is printed after it. If the function is a member of a cycle, the cycle number is printed between the function's name and the index number.

For the function's parents, the fields have the following meanings:

self	This is the amount of time that was propagated directly from the function into this parent.
children	This is the amount of time that was propagated from the function's children into this parent.
called	This is the number of times this parent called the function '/' the total number of times the function was called. Recursive calls to the function are not included in the number after the '/'.
name	This is the name of the parent. The parent's index number is printed after it. If the parent is a member of a cycle, the cycle number is printed between the name and the index number.

If the parents of the function cannot be determined, the word '<spontaneous>' is printed in the 'name' field, and all the other fields are blank.

For the function's children, the fields have the following meanings:

self	This is the amount of time that was propagated directly from the child into the function.
children	This is the amount of time that was propagated from the child's children to the function.

called This is the number of times the function called
 this child '/' the total number of times the child
 was called. Recursive calls by the child are not
 listed in the number after the '/'.

name This is the name of the child. The child's index
 number is printed after it. If the child is a
 member of a cycle, the cycle number is printed
 between the name and the index number.

If there are any cycles (circles) in the call graph, there is an
entry for the cycle-as-a-whole. This entry shows who called the
cycle (as parents) and the members of the cycle (as children.)
The '+' recursive calls entry shows the number of function calls that
were internal to the cycle, and the calls entry for each member shows,
for that member, how many times it was called from other members of
the cycle.

Copyright (C) 2012-2020 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification,
are permitted in any medium without royalty provided the copyright
notice and this notice are preserved.

Index by function name

```
[13] __GLOBAL__sub_I__ZN4GridC2Eii [5] seconds() [16]
LaplaceSolver::LaplaceSolver(Grid*)
[12] BC(double, double) [15] LaplaceSolver::initialize() [17]
LaplaceSolver::~~LaplaceSolver()
[4] SQR(double const&) [3] LaplaceSolver::solve(int, double) [18]
Grid::setBCFunc(double (*)(double, double))
[14] __static_initialization_and_destruction_0(int, int) [2]
LaplaceSolver::timeStep(double) [19] Grid::Grid(int, int)
```